

This document includes solution hints, but **NOT** necessarily the full solution!

Part I

Evaluation Page Part B

Name: _____ Signature: _____

1 Evaluation

This part of the exam has 152 questions, with a total of 416 points and 4 bonus points.

Part	Max. Points	Scored Points
A	60	
B	180	
	Total	

2 Evaluation Part B

Page	Points	Bonus Points	Score
5	19	1	
6	2	0	
7	18	0	
8	13	0	
9	10	0	
10	16	0	
11	12	0	
12	19	0	
13	2	0	
14	2	1	
15	10	0	
16	11	0	
17	7	0	
18	10	0	
19	5	0	
20	14	0	
21	5	0	
22	4	0	
23	2	0	
Total:	181	2	

Page	Points	Bonus Points	Score
24	5	2	
25	10	0	
26	4	0	
27	5	0	
28	14	0	
29	16	0	
30	10	0	
31	6	0	
32	10	0	
33	9	0	
34	9	0	
35	6	0	
36	12	0	
37	10	0	
38	8	0	
39	9	0	
40	9	0	
41	3	0	
42	12	0	
43	3	0	
44	9	0	
45	13	0	
46	3	0	
47	37	0	
48	3	0	
Total:	235	2	

Part II

Rules

Answer the questions within the space provided. If you do not have enough space, you can use the backside of the sheet. In that case clearly indicate that your answer continues on the backside.

3 Supporting Materials

This is an examination in writing, without the usage of any electronic devices, except a scientific pocket calculator. No restriction on the model of calculator that may be used, but no device with communication capability shall be accepted as a calculator. All other electronic devices are prohibited. Writing paper is available, writing instruments (pencil, pens, etc) have to be organized by the student.

- **Part A:** Without any supporting material, with calculator.
- **Part B:** With a self written summary (format A4, 8 sheets or up to 16 pages), with calculator

4 Procedure

1. Duration:
Part A: 1 hour = 60 minutes = 60 points.
(short break)
Part B: 3 hours = 180 minutes = 180 points.
2. Sign the first page in the provided space. With this you certify that you are only using permitted support material and you are complying to the rules.
3. Write your name on any detached or additional paper sheets. Sheets without a name will not be evaluated.
4. Use the provided paper for your solutions. Use the provided space in the forms and tables. If needed use scratch paper. Document your way to your solution as appropriate.
5. Each question has a defined number of maximum points associated.
6. If a question is unclear, make reasonable assumptions. Document your assumptions and provide a rationale.
7. Write clearly and legibly. Unclear or multiple solutions will not be evaluated.
8. There is a short break between part A and B. You have to sign into a list for a needed break during the examination parts. Only one person can leave the room for a short time.
9. If something is unclear, ask your supervisor in the room.

5 Time Management

Read first all questions. Make sure you distribute your available time to all the questions. To reduce disturbance, ask questions in the first 15 minutes of the exam period.

6 Multiple-Choice Questions

1. Try to answer all questions if possible. If you are not sure, choose the answer which seems the best one.
2. For the questions of type ○: Choose **exactly one** option with ⊗ (or √), which you think is the best match. With a correct answer you get the given number of points for that question.
3. For the questions of type ±: After a question or possibly incomplete sentence there are four answers or extensions. Evaluate each of them if they are true or false and mark them accordingly with '+' (true) or '-' (false). Independent if the question is formulated grammatically in singular or plural, it is possible that **0, 1, 2, 3, 4** of the choices are true. For three correct answers out of four you receive half of the points.
4. Wrong answers will have no penalty. Each question which has no answer is treated like a wrong answers and will be evaluated with zero points.
5. If you are changing your mind: cross out your old answer and clearly mark which answer is the new one.

May Dilbert be with you! ☺

Question 1 1 Point (Bonus)

This could be your own bonus question ☺.

1. What would be the answer?

Question 2 Points: [2]

In version control systems you are asked to provide a 'commit message': explain why you should provide such a message:

Solution: It allows to list them with the commits, so other users have a high level understanding of the commit.

Question 3 Points: [2]

In our drivers we implemented `Deinit()` functions, for example for the shell:

```
void SHELL_Deinit(void);
```

Explain the purpose of such `Deinit()` functions:

Solution: Used to de-initialize the driver, free up memory or semaphores/etc. Get the driver back to the reset state.

Question 4 Points: [5]

- (a) The boards we used (Robot and FRDM) have different number of LEDs available. Describe your strategy how you can write a common LED driver for all these boards (supporting multiple platforms with a single driver):

[3]

Solution: Using a platform define with macros to describe how many LEDs they have, and a macro to enable/disable it.

- (b) Describe with your approach how you can deal with a board which has no LED's at all:

[2]

Solution: Use a macro to disable it completely, or set number of LED to zero.

Question 5 Points: [10]

Given the following interfaces:

```
bool KEY_Get(void); /* return FALSE if key is pressed, TRUE otherwise */
void WAIT_Waitms(uint16_t ms); /* realtime waiting for the given
    milliseconds */

typedef enum {
    NO_KEY,      /* no key pressed */
    SHORT_KEY,   /* short key press */
    LONG_KEY     /* long key press */
} KEY_State;

KEY_State GetKey(void);
```

Implement the function `GetKey()` without the usage of interrupts, with following requirements:

1. Return `NO_KEY` if the key is not pressed.
2. If the key is pressed, debounce it for 50 ms.
3. If the key is pressed for less or equal than 500 ms, return `SHORT_KEY`.
4. For a short key press or long key press detection, the function shall not block longer than needed.
5. If the key is pressed for more than 500 ms, return `LONG_KEY`.

Solution:

```
PTADD = 0x00; /* set port as input */
PTAPE = 0xFF; /* enable pull ups */
while (PTAD!=0xFF); /* wait until port pull ups are working */
EnableInterrupts;
```

Solution:

```
KEY_State GetKey(void) {
    if (!KEY_Get()) { /* press */
        WAIT1_Waitms(50); /* simple debounce */
        if (!KEY_Get()) { /* still pressed? */
            cnt = 0;
            while (!KEY_Get() && cnt<=500) {
                WAIT_Waitms(1);
                cnt++;
            } /* wait until released */
            if (cnt>=500) { /* long press */
                return LONG_KEY;
            } else {
                return SHORT_KEY;
            }
        }
    } else {
        return NO_KEY;
    }
}
```

Question 6.....Points: [8]

You are using GIT as VCS. Explain what the following actions mean. Use '*remote repository*', '*local repository*', '*index*' and '*working directory*' in your answers.

(a) clone

[2]

Solution: Copy the remote repository as local repository.

(b) add

[2]

Solution: Add a new file to the index.

(c) push/sync

[2]

Solution: Add my changes from the index/local repository to the global repository.

(d) commit

[2]

Solution: Add my changes from the working directory/index to the local repository.

Question 7.....Points: [3]

List three different clock sources which can be used for the microcontroller we used in INTRO:

Solution: External crystal, external oscillator, internal clock.

Question 8.....Points: [3]

We used in our shell parser the following `xatoi()` interface:

```
byte UTIL1_xatoi(const unsigned char **str, long *res);
```

Explain why it is using `**str` and *not* `*str`:

Solution: It advances the parsing pointer, so `xatoi()` can be used sequentially on the input string.

Question 9.....Points: [3]

In our labs we used the RTOS with a tick timer period of 10 ms. List advantages and disadvantages for changing it to a period of 1 ms:

Solution: Pros: higher accuracy for wait/timing/scheduling (1). Cons: higher interrupt load (1). Performance tick counter needs to be changed too (1).

Question 10.....Points: [3]

Below is an extract of a header file:

```
/* drv.h */
#ifdef DRV_H_
#define DRV_H_
/* header file content follows here... */
#endif
```

Using the above header file, you find out that things are not working properly. Identify the problem.

Solution: Typo, should be `#ifndef`.

Question 11.....Points: [3]

List advantages and disadvantages of using Processor Expert for a project we did in INTRO:

Solution: Pros: hardware abstraction, provided libraries. Cons: learning phase, code not that optimized.

Question 12.....Points: [3]

The microcontroller we used in INTRO are using a table for the vectors: The vector table contains the address of the vector function. Explain what will happen during power-on reset if the whole vector table would be filled with zero (0x00) bytes:

Solution: Power-On reset will not work (no stack defined, no startup vector) (1.5). Code will divert to address 0x0, potentially crash. (1.5)

Question 13.....Points: [3]

Explain the fundamental differences between using `EnterCritical()` ... `ExitCritical()` and `xSemaphoreTake()` ... `xSemaphoreGive()` with respect to interrupts:

Solution: `EnterCritical` disables the interrupts (1), while the other method is using semaphores and itself does not disable interrupts (2).

Question 14.....Points: [4]

On the FRDM board we used an interrupt to detect button presses on the joystick shield.

(a) Why do we not know *which* button was pressed from the interrupt source? [1]

Solution: There is only one interrupt source for multiple interrupt pins (only one port interrupt, but multiple port pins).

(b) How can you know which button is pressed? [1]

Solution: Read the port pin value.

(c) We get an interrupt for a button. But when we check the pin status in the debouncing state machine we see that this pin is not pressed. What could be the most likely problem? [2]

Solution: If at the interrupt execution time the switch state already has been closed (port does not reflect the state at the interrupt time).

Question 15.....Points: [4]

Figure 1 shows three different circuits how to connect a switch to a microcontroller. Rate them either 'good', 'fair' and 'poor' with a short explanation.




Figure 1: Three Switches

Solution: (good, fair, poor), Current, defined state, internal pullup.
see <http://www.freewebs.com/maheshwankhede/ports.html>.

Question 16.....Points: [3]

The following `#define` might cause a problem using it:

```
#define INCREMENT(a)    a+1
```

Explain the problem and provide a solution how to fix it:

Solution: Textual replacement, order of operation might depend on usage. Better to use `((a)+1)`.

Question 17.....Points: [3]

Explain why FreeRTOS is using 'number of addressable units' as parameters to create a task, and not 'number of bytes':

Solution: Because stack needs to be aligned anyway on addressable units.

Question 18.....Points: [4]

An RTOS is using following interface for a 'wait' API call:

```
void vTaskDelay (portTickType xTicksToDelay);
```

- (a) Explain why FreeRTOS is using 'number ticks' as parameters to the delay routines and not 'number of milliseconds': [2]

Solution: Because the tick timer can be in smaller units than milli-seconds.

- (b) You need to wait for 200 milliseconds with `vTaskDelay()`, but you do not want to depend on the frequency of the tick timer. How can you do this? [2]

Solution: Using something like `vTaskDelay(200/portTICK_RATE_MS);`

Question 19.....Points: [3]

Evaluate the following statements about interrupt system for the FRDM (ARM Cortex M0+ core).

- ☐ \pm I cannot change the interrupt priorities of the Cortex M0+ interrupts with positive interrupt vector numbers.
- ☐ \pm Multiple interrupts can have the same priority on the Cortex-M0+.
- ☒ \pm I cannot change the interrupt priority of the negative interrupt vector numbers.
- ☒ \pm The ARM-Cortex M0+ core has only main priorities and no sub-priorities.

Question 20.....Points: [2]

List two advantages of using an optical quadrature encoder compared to a mechanical one:

Solution: Optical encoders do not have the problem of bouncing, and they can operate at faster signal frequency as there are no mechanical parts involved. Additionally they are free of maintenance or wear-out.

Question 21.....Points: [2]

List the 4 specific requirement for a realtime system:

Solution: The (a) right answer at the (b) right time, (c) independent of the system load, (c) in deterministic way.

Question 22.....Points: [3]

Explain the fundamental difference between a hard and a soft realtime system:

Solution: A hard realtime system has to comply with a hard time boundary: missing that boundary means failure. A soft realtime system might miss that time boundary, and the system still would be considered working properly, but missing the time boundary would mean a system degradation (not a failure).

Question 23.....Points: [2]

Explain the advantages and disadvantages of using Processor Expert for a system we have developed in INTRO:

Solution: It allows to abstract from the hardware, and to focus on the higher level of the application. It generates the driver code and saves development time. On the other side it necessary to learn a tool, and to configure it properly.

Question 24.....**Points: [2]**

Explain reasons why you would use a VCS:

Solution: Backup, restore, synchronization, tracking, short/long term undo, ownership, sandbox, branching, merging.

Question 25.....**Points: [2]**

Explain reasons when you should use a VCS:

Solution: Every time you need control over your code and/or when you work together with other persons on the same code.

Question 26.....**Points: [2]**

Explain disadvantages and advantages using macros:

Solution: Cons: Code might be harder to read (additional textual replacement), there are traps and pitfalls using macros. Pros: Easy way to configure code at compilation time (versus run time), better/faster code, helping the compiler to optimize code generation.

Question 27.....**Points: [2]**

Explain why synchronization between systems is needed:

Solution: Systems need to get aligned and synchronized with the real world. The real world time is continuous, while the computer time is clock based. To have computers attached to the real world, they need to synchronize with the time of the real world.

Question 28.....**Points: [2]**

List three different synchronization methods:

Solution: Interrupt synchronization, realtime synchronization, gadfly/polling synchronization.

Question 29.....**Points: [2]**

Explain a case where a synchronization between two systems does not make any sense:

Solution: If one system is so slow that it cannot react to the events from the other system. The system is overloaded and has no time to synchronize anyway.

Question 30.....Points: [2]

Explain reasons why we used a Shell in INTRO:

Solution: To communicate (i.e. inspect, configure, debug) our target.

Question 31.....Points: [2]

List reasons why you would use polling instead of interrupts for checking the state of a key:

Solution: Hardware does not support interrupts on that pin(s), no problem with missing a key press, plenty of computation time available.

Question 32.....Points: [2]

List key benefits of using open source operating system like FreeRTOS:

Solution: Sources available, free of charge, community working on development.

Question 33.....Points: [2]

Explain the difference between a pre-emptive and a cooperative multitasking:

Solution: Preemptive: scheduler is in charge who is executing, tasks get interrupted. In cooperative mode the task needs to return the control to kernel.

Question 34.....Points: [2]

What are stdin, stdout and stderr?

Solution: These are callbacks and standard handles (or pipes) used by the shell. Similar to what exists in Linux.

Question 35.....Points: [9]

FreeRTOS is used in priority based preemptive mode. Tick timer is set up for 10 ms, and the processor is running at maximum speed. Task T3 has priority 3 (highest priority), task T2 has priority 2 and task T1 has priority 1. The tasks have been created with `xTaskCreate()` before time $t=0$ ms, and the scheduler is started with `vTaskStartScheduler()` at the time $t=0$ ms. Draw in Figure 2 a timing diagram for the execution of the two tasks for the first 100 ms (after $t_0 = 0$ ms). Use a bar to indicate when a task is running. You can ignore the overhead in the task loop and the overhead in the RTOS/scheduler itself.

```

static portTASK_FUNCTION(T3, pvParameters) { /* priority 3 task */
    portTickType xLastTime = xTaskGetTickCount();
    for (;;) { /* task time is 5 ms including overhead */
        DoWorkFor5ms(); /* this needs 5 ms */
        vTaskDelayUntil(&xLastTime, 25/portTICK_RATE_MS);
    } /* loop forever */
}

static portTASK_FUNCTION(T2, pvParameters) { /* priority 2 task */
    for (;;) { /* task time is 10 ms including overhead */
        DoWorkFor10ms(); /* this takes 10 ms */
        vTaskDelay(30/portTICK_RATE_MS);
    } /* loop forever */
}

static portTASK_FUNCTION(T1, pvParameters) { /* priority 1 task */
    for (;;) { /* task time is 2 ms including overhead */
        DoWorkFor2ms(); /* this needs 2 ms */
        vTaskDelay(20/portTICK_RATE_MS);
    } /* loop forever */
}

```




Figure 2: Task Timing

Question 36.....Points: [2]

Evaluate following statements for the ANSI `strcmp()` and `strncmp()` library functions as used in the INTRO program:

- \ominus \pm `strcmp()` has one additional parameter compared to `strncmp()`.
- \oplus \pm `strcmp()` and `strncmp()` return both zero, if the two argument strings are equal.
- \oplus \pm Using `strncmp()` allow to compare two strings up to a given offset.

- ⊖ ± The parameter passing for `strncmp()` is usually more efficient than the parameter passing for `strcmp()`.

Question 37 **1 Point (Bonus)**

List a compelling reason why normally only debounced switches should be used at an interrupt generating input pin only:

Solution: As the switch is bouncing, it might create many (too many) interrupts on that input line. This will increase the system load and it might cause signaling multiple switch presses even if there was just one.

Question 38 **Points: [2]**

Determine the value of variable `j` after calling the function `foo()`:

```
#define SUM(a,b) (a)+b
#define MUL(a,b) a*b
#define CALL(a)  bar(a)

int bar(int i) {
    return MUL(i,SUM(3,MUL(5,10)));
}
int j;

void foo(void) {
    j = CALL(SUM(10,2));
}
```

38. (12*(3)+5*10) = 86

Question 39 **Points: [10]**

Consider following semaphore definition:

A semaphore `S` denotes an integer variable, which is accessed through the following two atomic operations:

```
#define DOWN(S)    while (S<=0); /* do nothing */ \
                  S = S-1;
#define UP(S)      S = S+1;
```

Using the two operations, it is possible for a program to implement a critical section or a mutex, as in the following example:

```
Semaphore S = 1; /* initialize semaphore to 1 */
```

Each process or task is able to protect its critical section:

```

while (1) {
    DOWN(S);
    /* in critical section */
    UP(S);
}

```

- (a) Is the implementation with UP() and DOWN() using a 'busy waiting'? Justify your answer. [2]

Solution: This semaphore definition involves the busy waiting because a waiting process is looping for the shared variable S.

- (b) The tasks are using a priority based scheduling. Is it possible that the above solution leads to a priority inversion problem? Provide an rationale for your answer and provide an example why. [2]

Solution: Yes, the above solution can lead to the priority inversion problem if a high-priority process is waiting for a low-priority process to leave the critical section and signal the shared variable S.

- (c) Implement in pseudo code the tasks A, B, C, D and E. Use UP() and DOWN() on the needed semaphores to implement following scenario: After both A and B have finished their work, task C can start. After task C has finish its work, both D and E can start. [6]

Solution: The relations among processes can be represented in the following diagrams.

```

A \    / D
S1 \  /
    C S3
S2 /  \
B /    \ E

```

The program can be as follows:

```
semaphore S1, S2, S3 = 0, 0, 0;
```

Process A: (2p)

- do work of A

```
UP(S1); /* Let C start */
```

Process B: (2p)

- do work of B

```
UP(S2); /* Let C start */
```

Process C: (4p)

```

DOWN(S1); /* Block until A finished */
DOWN(S2); /* Block until B finished */
- do work of C
UP(S3); /* Let D or E start */
UP(S3); /* Let D or E start */

```

Process D: (2p)

```

DOWN(S3); /* Block until C finished */
- do work of D

```

Process E: (2p)

```

DOWN(S3); /* Block until C finished */
- do work of E

```

Total: 10

Question 40.....Points: [4]

Explain the *Priority Inversion* problem and illustrate it with an example:

Solution: A priority inversion problem is a situation in which a low-priority process is blocking a high priority process.

Question 41.....Points: [2]

We are using a tick timer with a period of 10ms a preemptive RTOS (FreeRTOS). This tick timer is as well used for the Trigger module as implemented in INTRO. Describe the consequences for the application and RTOS if this tick timer period gets changed from 10 ms to 200 ms, with the example of our INTRO application at the end of the semester.

Solution: A preemptive scheduling would happen only every 200 ms. Triggers would have the resolution of 200 ms, and this would impact our debouncing (it would take longer). Wait/Waituntil would have a resolution of 200 ms.

Question 42.....Points: [10]

A device is connected to a microcontroller as in Figure 3. It is using a data bus (PTB: Data) and two control signals (PTA0: status; PTA1: control). The protocol to send an 8bit data byte is illustrated in Figure 3. Implement the protocol using a *Gadfly* synchronization method.

Note: use pseudo code as SetPA0asInput(), but make sure with comments and naming to make clear what it does

- (a) Implement the function `Init()` which initializes the ports and any other required variables.

[5]

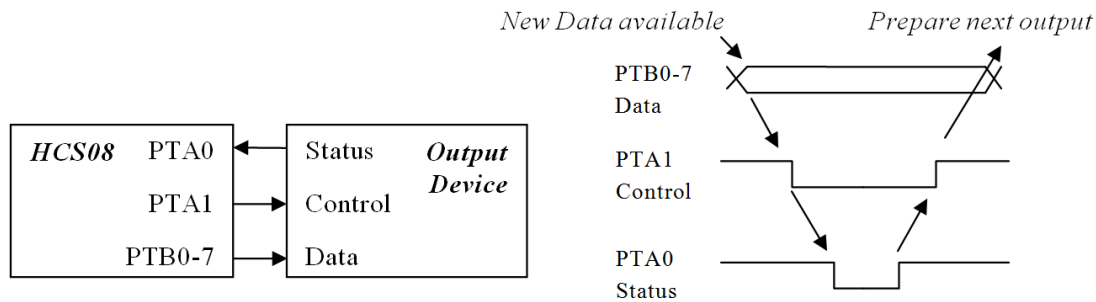


Figure 3: Output Device

Solution:

```

void Init(void) { (3p)
    PTAD_A0 = 1; /* default pin value to 1 */
    SetPTA0asInput; /* pin as input */

    PTAD_A1 = 1; /* pin set to 1 */
    SetPTA1asOutput;

    PTBD = 0; /* init port */
    SetPTBasOutput; /* PTB as output */
}

```

- (b) Implement the function **Send(char)** which transmits a single byte.

[5]

Solution:

```

void Send(char ch) {
    /* optional: wait until PTA0 is high */
    PTBD = ch;
    PTAD_A1 = 0; /* signal to device */
    while (PTAD_A0==1); /* wait until signal goes down */
    while (PTAD_A0==0); /* wait until signal goes up */
    PTAD_A1 = 1; /* signal finish transfer */
}

```

Total: 10

Question 43.....Points: [4]

- (a) Determine for the 8bit (binary reflected) Gray code **0x47_g** the corresponding 8bit Binary code_b:

[2]

(a) 122/0x7A_b**Solution:** Gray to Binary:0x47_g, then use MSB, then add each next bit (ignore carry):0100'0111 → 0111'1010 → 0x7A_b

- (b) Determine for the 8bit (binary reflected) Gray code $0x74_g$ Code the corresponding 8bit Binary code_b: [2]

(b) 88/0x58

Solution: Gray to Binary:

$0x74_g$, then use MSB, then add each next bit (ignore carry):

$0111'0100 \rightarrow 0101'1000 \rightarrow 0x58_b$

Total: 4

Question 44.....Points: [6]

Consider following ANSI-C implementation:

```
typedef enum LED_Set {
    LED_0 = 1,
    LED_1 = 2,
    LED_2 = 4,
    LED_3 = 8
} LED_Set;
```

Alternatively it can be written as:

```
#define LED_0 1
#define LED_1 2
#define LED_2 4
#define LED_3 8
```

- (a) List pros and cons for using enumerations for above example: [2]

Solution: Pros: symbolic debugging, using a type. Cons: it is using an int (which might be an overkill). Enums are int's in C.

- (b) List pros and cons for using #define's for above example: [2]

Solution: Pros: normal int type, is portable, can calculate with int's (masking, binary operations). Con: no symbolic debugging, no type for memory allocation, no type

- (c) Explain the reason why the values 1, 2, 4, 8 are used (and not 0, 1, 2, 3): [2]

Solution: Possible to use the LED's as bit set (efficient for passing multiple LED's in one argument).

Question 45.....Points: [2]

Which three critical hardware settings do you have to configure on the FRDM board to communicate using an RS-232-to-USB connection. Use the shell we used in INTRO as example.

Solution: Port/SCI to use, Baudrate, protocol (8N1).

Question 46.....Points: [5]

In this example FreeRTOS is used in priority based pre-emptive mode. Tick timer is set up to 10 ms, and the processor is running at maximum speed. The two LED's are switched off at the beginning. Task T1 has priority 3, and task T2 has priority 2. T1 has to work for 15 ms, and T2 for 2 ms (the overhead within the for loop can be ignored). Both tasks have been created with `xTaskCreate()` right after each other, and the scheduler has been started with `vTaskStartScheduler()` at the time `t=0` ms.

```
static portTASK_FUNCTION(T1, pvParameters) { /* priority 3 task */
    portTickType xLastTime = xTaskGetTickCount();
    LED1_Off();
    for (;;) { /* task time is 15 ms including overhead */
        DoWorkFor15ms();
        LED1_Neg();
        vTaskDelayUntil(&xLastTime, 25/portTICK_RATE_MS);
    } /* loop forever */
}

static portTASK_FUNCTION(T2, pvParameters) { /* priority 2 task */
    portTickType xLastTime = xTaskGetTickCount();
    LED2_Off();
    for (;;) { /* task time is 2 ms including overhead */
        LED2_Neg();
        DoWorkFor2ms();
        vTaskDelayUntil(&xLastTime, 30/portTICK_RATE_MS);
    } /* loop forever */
}
```

Draw in Figure 4 a timing diagram for the execution of the two tasks for the first 100 ms (after `t0=0ms`). Indicate the state of both LED's (LED1 and LED2). Use a bar to indicate when a task is running, and a bar to indicate the time when a LED is on. You can ignore the overhead in the task loop and the overhead in the RTOS/scheduler itself.

Solution: Keep in mind that the tick task counter resolution is 10ms: as such the `vTaskDelayUntil()` will take the current time at 10 ms boundary (e.g if the 'real' time is at 15ms, this corresponds to a tick counter of 1 (at 10 ms). If you make now a `vTaskDelayUntil()`, the time is taken from the tick counter and added. So if T2 calls `vTaskDelayUntil()` at the real time of 17 ms, this corresponds to a tick time of 10ms. Adding 30ms means that the task will be ready at 40 ms (however only, if not a higher priority task is running).

- correct LED toggling
- correct task length (15ms and 2ms)
- correct DelayUntil()




Figure 4: Task Timing

Question 47.....Points: [2]

Explain the difference between `xQueueReceive()` and `xQueuePeek()` in FreeRTOS:

Solution: `xQueueReceive()` removes the element from the queue. `xQueuePeek()` only checks if the queue contains an element, but does not remove it.

Question 48.....Points: [12]

Given a system using *Priority Ceiling* and three tasks (J_1 (Priority 1, highest priority), J_2 (priority 2) and J_3 (priority 3)) and two semaphore S_1 and S_2 . The tasks are getting a lock of a semaphore with LS_x , and return the lock with US_x . The tasks are using the semaphores according to following timing and sequence:

- $J_1 = \{10\text{ms}, LS_1, 10\text{ms}, US_1, 10\text{ms}\}$: total 30ms task time.
- $J_2 = \{10\text{ms}, LS_2, 10\text{ms}, LS_1, 10\text{ms}, US_1, 10\text{ms}, US_2, 10\text{ms}\}$: total 50ms task time.
- $J_3 = \{10\text{ms}, LS_1, 30\text{ms}, US_1, 10\text{ms}\}$: total 50ms task time.

(a) Calculate the *Priority Ceiling* for S_1 : [2]

(a) $J_1 \rightarrow 1$

(b) Calculate the *Priority Ceiling* for S_2 : [2]

(b) $J_2 \rightarrow 2$

(c) The delta between t_n and t_{n+1} in Figure 5 is 10 ms. Following times are defined: [8]

- J_3 starts at time t_0 .
- J_2 starts at time t_2 .
- J_1 starts at time t_4 .

Complete and extend the timing diagram in Figure 5 indicating the task execution times. Mark with an LS_x the request for a semaphore x , and with US_x the release of a semaphore.

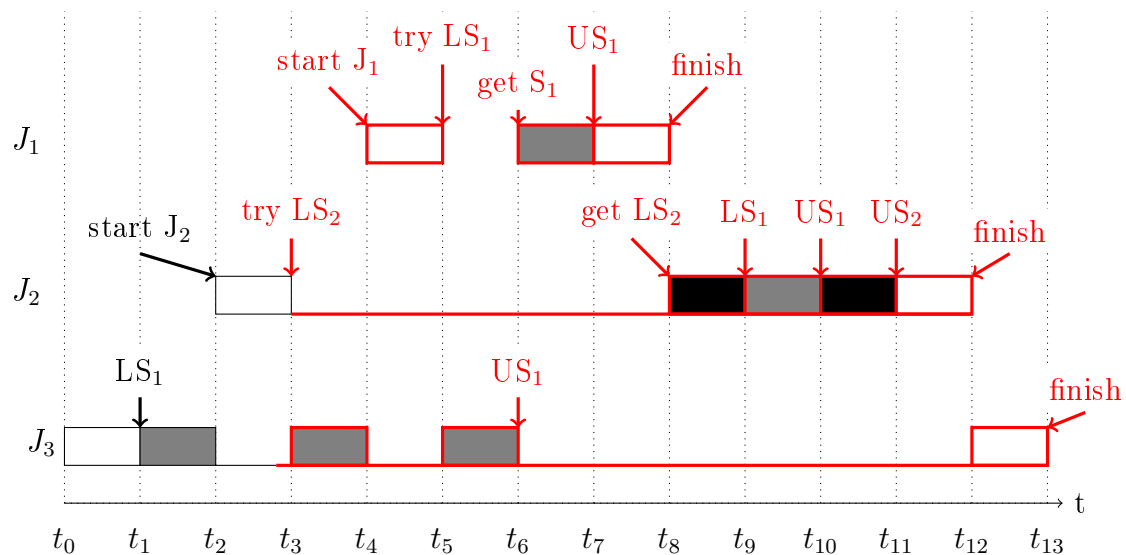


Figure 5: PCP

Question 49.....Points: [2]

The PWM signal used to drive the INTRO motor is 'low active'. Explain what this means.

Solution: The speed of the motor is proportional to the low level duty cycle of the signal: the more the signal is low, the higher the motor voltage will be.

Question 50.....Points: [1]

The following code

```
s = k0*e + k1*(e-eprev)/Ta;
eprev = e;
```

Implements in software a closed loop control using the following method:

☒ PD

☐ ID

☐ PI

☐ P

☐ PID

Question 51.....Points: [2]

Consider the H-Bridge we used for our INTRO lab. You have to implement a software which is able to drive the motor, from -100% (full speed backward) to +100% (full speed forward). Explain how you realize this with the signals you have available.

Solution: The range from 0 to 100% is realized with the PWM duty cycle (which is low active on our platform). 0 will be no speed, and 100% will be full speed. The direction signal (DIR) is used to drive the motor either forward (e.g. HIGH) or backward (e.g. LOW).

Question 52.....Points: [2]

- (a) Determine for the 8bit **Binary Code 0x31_b** the corresponding 8bit (binary reflected) Gray code_g: [1]

(a) 0x29_b

Solution: Binary to Gray:

0x31 EXOR (0x31»1)

= 0x31 EXOR 0x18 = **0x29_b** (or **0b0010'1001** or **41_d**)

- (b) Determine for the 8bit (Binary Reflected) **Gray Code 0x8F_g** the corresponding 8bit Binary Code_b: [1]

(b) 0xF5_b

Solution: Gray to Binary:

0x8F_g, then use MSB, then add each next bit (ignore carry):

0xF5_b (or **0b1111'0101** or **245_d**)

Question 53.....Points: [1]

Given the pseudo code for a closed loop control system:

```
esum += e;
s = a*e + b*Ta*esum;
```

This software controller implements a

- ☐ PD controller
☐ ID controller
☒ **PI controller**
☐ P Controller
☐ PID Controller

Question 54.....Points: [1]

For the digital encoder (using C1 and C2 signals) and the motor (as used in the INTRO lab) following applies:

- ☐ ± The quadrature encoder produces with C1 and C2 a 4 bit Gray encoded signal.
☐ ± In case C1 is not working or not available (only C2 is producing a signal), then we still can determine the motor direction.

- ⊖ ± With minimal four consecutive codes it is possible to determine the motor direction.
- ⊖ ± Using the two signals C1 and C2 it is possible to determine the absolute motor position.

Question 55.....Points: [2]

Given following program to debounce a switch on a board. `RawKeyPressed()` returns 0 for if the switch is pressed. The function `DebounceSwitch()` gets called periodically every 3 ms.

```
bool DebounceSwitch(void) {
    static word16 State = 0;
    State = (State<<1) | !RawKeyPressed() | 0xF800;
    if (State==0xFC00) return TRUE;
    return FALSE;
}
```

Determine the switch release debouncing time in milliseconds which is used to debounce the switch.

55. 30 ms, or number of zero bits (1111'1100'0000)

Solution:

```
1111'1000'0000'0000 FALSE (key not pressed)
1111'1000'0000'0001 FALSE (key pressed: !0 => 1 shifted in)
1111'1000'0000'0011 FALSE
1111'1000'0000'0111 FALSE
...
1111'1111'1111'1111 FALSE (key pressed stable for a long time)
Release (input 0)
1111'1101'1111'1110 FALSE
...
1111'1101'1010'0000 FALSE
...
1111'1010'0000'0000 FALSE
1111'1100'0000'0000 TRUE
```

Question 56.....Points: [3]

Processor Expert components are using the concept of *Method*, *Property* and *Event*. Give examples of this concept using the FreeRTOS component.

- (a) 3 FreeRTOS examples for *Methods*:

[1]

Solution: vTaskCreate(), vTaskStartScheduler(), vWaitUntil()

- (b) 3 FreeRTOS examples for *Properties*:

[1]

Solution: InitialStackSize, Taskname length, preemptive or non-preemptive

- (c) 3 FreeRTOS examples for *Events*:

[1]

Solution: onStackOverflow(), onTimerTick(), MemoryFailed()

Question 57.....Points: [1]

You connect a push button which is *not* debounced to interrupt input pin, using an external pull-up resistor. You configure the pin to generate an interrupt whenever there is a falling edge. What are the things you have to consider?

Solution: You have to disable immediately (as fast as possible) the interrupts for the input. Then you need to properly debounce the switch. Be prepared that you might have an interrupt flag pending, so you need to acknowledge the interrupt at the debouncing.

Question 58.....Points: [1]

Explain why there is the word 'quadrature' used to denote a quadrature encoder/decoder:

Solution: Because the signals consist of two sinus-like signals, which are 90 degrees out of phase (or a quadrant of a full circle).

Question 59 2 Points (Bonus)

You are implementing a state machine which performs a debouncing. You want to use the same state machine driver for different keys at the same, and depending on the application you want the state machine calling different things. What are you using to implement the *user/application defined* function calls which are evaluated at run time?

Solution: Function pointers.

Given the the FreeRTOS task scheduling pattern as shown in Figure 6 for the first 100 milliseconds. The RTOS is in preemptive scheduling mode. The RTOS IDLE task is running with priority `tskIDLE_PRIORITY`. The tasks T1 and T2 are implemented as following:


Question 60.....Points: [8]

Figure 6: FreeRTOS Task Schedule

```
static portTASK_FUNCTION(T1, pvParameters) {
    for (;;) {
        DoWorkT1();
        vTaskDelay(delayT1);
    } /* loop forever */
}

static portTASK_FUNCTION(T2, pvParameters) {
    for (;;) {
        DoWorkT2();
        vTaskDelay(delayT2);
    } /* loop forever */
}
```

(a) Determine the value of `delayT1` and `delayT2` in *milliseconds*: [2]

(a) delayT1: 20 ms, delayT2: 20 ms

(b) Each task does some work which is *not* using any blocking RTOS calls. Determine the time of work for `DoWorkT2()` and `DoWorkT1()` in *milliseconds*: [2]

(b) DoWorkT1: 5 ms, doWorkT2: 10 ms

(c) What is the tick timer period of the RTOS? Explain the reasoning behind your answer. [2]

Solution: Tick timer is 5 ms, as the delays are a multiple of it, and task switching happens from the IDLE task at 5 ms boundaries.

(d) Could it be possible that T1, T2 and IDLE share the same priority level? Justify your answer. [2]

Solution: Yes, this is possible, as within at least the first 100 ms there is no task running at the same time

Question 61.....Points: [2]

Given the following enumeration type:

```
typedef enum { RED, BLUE, GREEN, YELLOW, BLACK } ColorsEnum;
```

For the compiler we used in INTRO, the following applies for ANSI-C (with default compiler settings):

- ☐ sizeof(ColorsEnum)==1
- ☐ sizeof(ColorsEnum)==2
- ☒ sizeof(ColorsEnum)==4
- ☐ sizeof(ColorsEnum)==5
- ☒ sizeof(ColorsEnum)==sizeof(int)

Question 62.....Points: [4]

Given following declaration of a trigger:

```
#define TRG_LAST      4
typedef void (*TRG_Callback)(void*);

typedef struct TriggerDesc {
    uint16_t triggerTicks;
    TRG_Callback callback;
} TriggerDesc;

static TriggerDesc TriggerList[TRG_LAST];
```

The triggers are initialized with following function:

```
void Init(void) {
    static volatile uint8_t i;

    for(i=sizeof(TriggerList)/sizeof(TriggerDesc); i>0; i--) {
        TriggerList[i].triggerTicks = 0;
        TriggerList[i].callback = NULL;
    }
}
```

(a) Evaluate the following statements:

[2]

- ☐ \pm The function `Init()` will never finish.
- ☐ \pm Every element of `Triggerlist` is initialized with with the same values.
- ☐ \pm Using `static volatile` for the variable `i` ensures that the function `Init()` is reentrant.
- ☐ \pm Using `volatile` for the variable `i` ensures that the function `Init()` is reentrant.

(b) Using above implementation, you notice that your application is working in a strange way, it even crashes after a while. Explain the reason:

[2]

Solution: The 'for' loop will iterate through index 4 to 1. But the array has the index 3 to 0. As such, the loop will write somewhere in memory, and part of the triggers are not initialized.

Question 63.....**Points: [5]**

Given following implementation of an event module:

```
#define GET_EVENT(event) \
    (bool)(EVNT_Events[(event)/8]&(0x80>>((uint8_t)((event)%8)))
#define CLR_EVENT(event) \
    EVNT_Events[(event)/8] &= ~(0x80>>((uint8_t)((event)%8)))

static uint8_t EVNT_Events[((EVNT_NOF_EVENTS-1)/8)+1];

bool EVNT_GetEvent(EVNT_Handle event) {
    bool isSet;
    EnterCritical();
    isSet = GET_EVENT(event);
    ExitCritical();
    return isSet;
}

void EVNT_ClearEvent(EVNT_Handle event) {
    EnterCritical();
    CLR_EVENT(event);
    ExitCritical();
}

uint8_t EVNT_CheckEvents(void) {
    uint8_t i;

    EnterCritical();
    for(i=0; i<sizeof(EVNT_Events)/sizeof(EVNT_Events[0]); i++) {
        if (EVNT_GetEvent(i)) {
            EVNT_ClearEvent(i);
            break;
        }
    }
    ExitCritical();
    return i; /* return the event which was set */
}
```

- (a) Using that implementation and calling `EVNT_CheckEvents()`, you find out that your periodic interrupt timer does not work any more. Explain the problem and how to fix it: [3]

Solution: The `Enter/ExitCritical()` are nested, thus it will disable the interrupts always. The solution is to use inside the `for()` loop the `GET_EVENT()` and `CLR_EVENT()` macros.

- (b) After you have fixed the previous problem, the application still does not work correctly. Somehow `EVNT_CheckEvents()` does not properly return events which have been set. Explain the problem and how to fix it: [2]

Solution: The `for()` loop is not properly iterating through the bit array. The loop should use a test on `i < EVNT_NOF_EVENTS` instead.

Question 64.....Points: [2]

Evaluate following statements in the context the RTOS startup, as used in the INTRO lab (FreeRTOS):

- ⊕ ± All interrupts are disabled during the (ANSI) startup code.
- ⊕ ± Global interrupts are enabled at the entry point of `main()`.
- ⊕ ± Interrupts for RTOS timers are disabled at the entry point of `main()`.
- ⊕ ± Interrupts for RTOS timers get enabled during startup of the scheduler.

Question 65.....Points: [3]

Explain why it is necessary to implement a timeout using the delta-time approach for measuring the speed with a quadrature encoder:

Solution: If there is no movement or the movement is very slow, then there would be no state change. As such, a timeout needs to be implemented in order to limit the delta time value.

Question 66.....Points: [3]

If you can choose between the delta-time and delta-pos method for quadrature decoding for fast (many steps) state changes, which one would you prefer and why?

Solution: The delta-pos approach, as there would be sufficient position changes in a small amount of time.

Question 67.....Points: [3]

If you can choose between the delta-time and delta-pos method for quadrature decoding for slow (few steps) state changes, which one would you prefer and why?

Solution: The delta-time approach, measuring the time for one or few steps. There would not be enough or now step for a given (short) time to count.

Question 68.....Points: [3]

Explain why a tacho implementation as we have implemented is not an exact measurement, but rather a speed estimation?

Solution: The sensor values are always lagging (behind in time, backward looking), the encoder itself has errors, there could be a drift between wheels and ground, and there are quantization errors.

Question 69.....Points: [3]

Name a data structure which can be used to efficiently implement a delta-pos algorithm for estimating a tachometer.

Solution: A ring buffer can be used for this: new values would overwrite the oldest entries.

Question 70.....Points: [3]

What is the difference between a Mutex and a normal Semaphore in FreeRTOS?

Solution: The Mutex is a special case of a semaphore. A Mutex is used to grant exclusive access to a shared resource, while semaphores can be used for inter-process communication. The Mutex implements Priority Inheritance, which means that a process waiting on the Mutex inherits the priority of the process holding that Mutex. A Mutex has to be always returned, while for a semaphore this is not required.

Question 71.....Points: [3]

How many processes and priorities are needed to demonstrate a Priority Inversion?

Solution: At least three processes with different priorities are needed for a Priority Inversion. One process with low priority to hold the semaphore/resource, one process with medium priority which blocks the low priority process, and a process with high priority which waits for the semaphore/resource.

Question 72.....Points: [3]

Given the two Semaphores S1 and S2 and the three processes PL (high priority (priority 3)), PM (medium priority (priority 2)) and PL (low priority (priority 1)). Semaphore S1 is used by PL and PM, S2 is used by PL and PH. Determine the *Priority Ceiling* of S1 and S2.

Solution: S1 has a Priority Ceiling of 2, S2 has a Priority Ceiling of 3.

Question 73.....Points: [3]

If you use a Semaphore in FreeRTOS from an interrupt service routine (ISR), is there something special to consider?

Solution: To use RTOS API calls within an ISR, the API routines with a *fromISR* suffix has to be used, for example `xSemaphoreTakeFromISR()` and `xSemaphoreGiveFromISR()`.

Question 74.....Points: [1]

In which state is a process in FreeRTOS when it is waiting on a Semaphore?

Solution: It is in the *Blocked* state.

Question 75.....Points: [3]

Using a network stack like RNet, which things do you have to initialize?

Solution: Call the stack initialization routine (this initializes the low level driver), assign application message handler, assign own node address

Question 76.....Points: [2]

RNet uses four different layers in the stack. List the layers with their acronyms.

Solution: PHY (Physical Layer), MAC (Media Access Layer), NWK (Network Layer), APP (Application Layer)

Question 77.....Points: [1]

What is a CRC and for what it is used for?

Solution: CRC (Cyclic Redundancy Check) implements a check-sum to detect errors in data. In communication systems it is used to check the correctness of data packets transmitted.

Question 78.....Points: [1]

What is a so called *Ad-Hoc Network*?

Solution: A network with two or more nodes which does not need a centralized or fixed network service infrastructure. This differentiates it from other networks like a Wireless LAN with an access point.

Question 79.....Points: [2]

List pros and cons of an *Ad-Hoc Network*.

Solution: Pros: No central infrastructure required, network can be built 'on the fly' and on demand. No need for an access point or network coordinator which needs to run all the time. Cons: Resources and energy needed to run the access point or network coordinator. Network might not be scalable or suitable for large networks. No administration of the network.

Question 80.....Points: [1]

Using Segger JScope, which variables can be monitored by that tool without special extensions of the program itself?

Solution: Only global (external linkage) variables can be monitored. It is not possible to monitor local (stack) variables or CPU registers.

Question 81.....Points: [1]

Which communication way did we use with Segger JScope?

Solution: We used the debug protocol (SWD) for this, over the debug cable.

Question 82.....Points: [1]

What is the purpose of a radio transceiver?

Solution: A radio transceiver is used for over-the-air communication. The transceiver transforms the data packets suitable for physical communication channel, encodes and decodes them.

Question 83.....Points: [1]

Name three software components/layers of the RNet communication stack we used.

Solution: PHY (physical), MAC (media access), NWK (network), APP (application).

Question 84.....Points: [1]

What does *Payload Packaging* mean?

Solution: With Payload Packaging each layer of a network stack can add/remove additional data at the beginning and/or end of the payload of the higher/lower network layers.

Question 85.....Points: [1]

You need to drive a H-Bridge IC with a PWM signal. What fundamental parameter you have to check in the data sheet of the IC?

Solution: The maximum/minimum PWM frequency you can use. Each IC has its own specific range.

Question 86.....Points: [1]

How can the power dissipation of the H-Bridge for of a PWM controlled device be reduced?

Solution: In general, the data sheet should provide information about the power dissipation at given frequencies. Typically higher PWM frequencies will cause higher dissipation due the switching activity, so a frequency in MHz range does not make any sense. Many motor controllers use a frequency above 20 kHz for quieter operation.

Question 87.....Points: [1]

If a closed control loop controller is using a PWM as output signal, what should be considered?

Solution: Depending on the PWM hardware, the PWM duty should be changed as such that a current running PWM duty cycle is not affected. Ideally the PWM duty cycle is changed at the end of a PWM period.

Question 88.....Points: [1]

What is the common purpose of using a Queue or a Semaphore/Mutex?

Solution: To ensure that access to data is done in a controlled way without race conditions or undesired side effects.

Question 89.....Points: [1]

Lock() and Unlock() are functions used for which programming primitive or design pattern?

Solution: For Semaphores or Mutex.

Question 90.....Points: [1]

The speed of a motor is proportional to the voltage applied to the motor. So why is it necessary to use a closed control loop?

Solution: To deal with different load, mechanical differences and external influences.

Question 91.....Points: [2]

You have to implement a position controller (closed loop)? Which type do you choose: PD or PI?

Solution: PD. The system already has an integral part in it and integrating the position error would either be a very large amount or would need to be scaled down so largely that it does not make much sense. As such, the PD would be sufficient.

Question 92.....Points: [2]

For our Sumo robot, identify the following parts of a closed loop control system: Controller, Actuator, Sensor, System.

Solution: Controller: Microcontroller with PID control loop implementation. Actuator: H-Bridge. Sensor: quadrature decoder (speed/position), System: Sumo Robot motors

Question 93.....Points: [2]

What is the purpose of the *Anti-Windup*?

Solution: To limit the total sum and impact of the integral value.

Question 94.....Points: [2]

Why is it better to use 'industrial qualified' SD cards instead of using 'consumer' ones?

Solution: The industrial ones usually are more robust and have well defined parameters or data sheet information. Consumer type cards might very differ in access time, wear out and programming time.

Question 95.....Points: [1]

Provide an implementation way where you can store calibration data?

Solution: Calibration data can be stored for example in the microcontroller internal flash, either at compilation time (constants in code) or at runtime with programming the FLASH memory.

Question 96.....Points: [3]

What do you have to consider if you are programming the internal FLASH memory at run time?

Solution: FLASH is typically organized in pages and blocks, and programming is done on page or block level. You need to make sure that you separate you program from the data. Interrupts typically needs to be disabled during flash programming. Typically the program memory is not accessible during flash programming, so the flash programming 'applet' needs to run in RAM.

Question 97.....Points: [1]

Why did we use a digital interface to the reflectance sensor and not an analog one?

Solution: Digital pins were easily available, while analog pins are usually limited on a microcontroller.

Question 98.....Points: [1]

Even if all sensors of the reflectance sensor are exposed to the same conditions, the measurement will differ from sensor to sensor. Why?

Solution: The resistors and the capacitors used have their own tolerances, resulting in different measurements. This results into different discharging time.

Question 99.....Points: [1]

The reflectance sensor array has two red LED's. What's the purpose of it?

Solution: To visually indicate if the sensors are used.

Question 100.....Points: [1]

How could you modify the reflectance sensor array hardware to minimize the power consumption?

Solution: Removing the red LED's or the current limiting resistor for each red LED.

Question 101.....Points: [3]

List typical external errors for the reflectance sensor array we have used, and how the errors can be reduced or limited.

Solution: Sun light, room ambient light, light from other infrared LED's. Errors can be minimized with proper sensor placement and distance to the ground, using shielding, multiple reads of sensors (averaging, etc). Calibrate the sensors for the given environment.

Question 102.....Points: [3]

List several functions or macros which are used in FreeRTOS to control task execution.

Solution:

- `uxTaskPriorityGet/Set()`
- `vTaskSuspend/SuspendAll()`
- `vTaskResume()`, `vTaskResumeAll()`, `vTaskResumeFromISR()`,
- `taskYIELD()`
- `taskENABLE_INTERRUPTS()`, `taskDISABLE_INTERRUPTS()`
- `taskENTER_CRITICAL()`, `taskEXIT_CRITICAL()`
- `vTaskDelay()`, `vTaskDelayUntil()`

Question 103.....Points: [1]

If several tasks are pending, which task gets executed first by a priority-based preemptive scheduler?

Solution: The one with the highest priority.

Question 104.....Points: [1]

Why could be a call to `vTaskEndScheduler()` make sense?

Solution: If only for a certain time an RTOS is needed in an application. For example an application only needs to run multiple tasks and uses the RTOS services in a special mode, but otherwise the application does not want the runtime overhead (e.g. for going into a very deep low power mode). On the other hand, the application could suspend the scheduler too.

Question 105.....Points: [1]

List the five memory allocation schemes in FreeRTOS.

Solution:

1. Scheme 1: only allocation, no deallocation.
2. Scheme 2: blocks are deallocated, but not merged.
3. Scheme 3: wrapper to `malloc()` and `free()`.
4. Scheme 4: same as Scheme 3, but merges free blocks.
5. Scheme 5: supports multiple memory regions.

Question 106.....Points: [1]

What was the purpose of the Shell queue we have implemented in the INTRO application?

Solution: The queue is used to pass data from multiple tasks to the Shell task. That way the other tasks send the data through the queue, instead of directly accessing the UART resource. So the queue is used as synchronization design pattern.

Question 107.....Points: [1]

What are the two different categories of timeliness?

Solution: Absolute and relative.

Question 108.....Points: [1]

List three different clock sources for a microcontroller:

Solution: External crystal, external oscillator, internal reference clock.

Question 109.....Points: [1]

In a microcontroller you can have different kinds of clocks, list at least two:

Solution: Memory Bus clock, CPU clock, Peripheral clock (A/D, USB, ...).

Question 110.....Points: [1]

List a fundamental requirement for a realtime operating system:

Solution: Exact time base.

Question 111.....Points: [1]

List pros and cons of using an internal reference clock (compared to an external clock generator):

Solution: Pros: less parts or PCB components, simple, less expensive. Cons: less accurate, may need calibration, uses more power (not ideal for low power applications).

Question 112.....Points: [3]

List reasons why to use an RTOS:

Solution: Running multiple things easily in parallel, scalability of the application, using RTOS services (queues, semaphore, mutex, memory allocation, events, ...)

Question 113.....Points: [2]

What is the difference between preemptive and non-preemptive scheduling:

Solution: In a preemptive operating system the RTOS distributes the processing time, tasks get suspended by the RTOS. In a non-preemptive (cooperative) operating system the tasks are passing actively the control back to the operating system.

Question 114.....Points: [2]

List the three standard I/O channels used by the Shell:

Solution: stdin, stdout, stderr

Question 115.....Points: [1]

List the four fundamental states of a task in an operating system:

Solution: Run, Idle, Wait, Ready

Question 116.....Points: [1]

What does the abbreviation *HAL* mean?

Solution: Hardware Abstraction Layer: an interface or similar in software to abstract from the actual hardware used.

Question 117.....Points: [2]

What is the difference between two different vendors offering a microcontroller based on an ARM Cortex M0+?

Solution: The processor core is the same (instruction set, registers), although it can be configured by the vendor in some details (e.g. number of interrupt levels). But everything around the core (e.g peripherals like USB, CAN, UART, ...) are vendor specific.

Question 118.....Points: [2]

How many hardware stack pointers implements an ARM Cortex M0+ microcontroller?

Solution: Two (Main SP and Process SP).

Question 119.....Points: [2]

The ARM Cortex M0+ architecture has 16 32bit registers. Are they all freely usable by the application code?

Solution: No, only R0 - R12. Other registers are used by the hardware R13 \Rightarrow SP, R14 \Rightarrow LR, R15 \Rightarrow PC R16 \Rightarrow xPSR

Question 120.....Points: [2]

Which Processor Expert generated macros can you use to protect a critical section?

Solution: EnterCritical() and ExitCritical()

Question 121.....Points: [1]

What could a reason to use the `volatile` qualifier for a loop variable?

Solution: To prevent compiler optimizations.

Question 122.....Points: [1]

What is the advantage of using *Gadfly Synchronization* compared to *Interrupt Synchronization*?

Solution: Less overhead with interrupt latency, simpler program flow, better if only to wait for a very short time.

Question 123.....Points: [1]

What does this mean: *an instruction is atomic*?

Solution: It means that an instruction is always executed as a whole, that means it cannot be interrupted and in between something different can be executed.

Question 124.....Points: [1]

What is the fundamental problem of 'shared data' in the context of interrupts?

Solution: An interrupt can happen any time. As such, if shared data is accessed from the main application and the interrupts, then access to the shared data could be interrupted too. As such, inconsistent data or race conditions could happen.

Question 125.....Points: [2]

The implementation of `EnterCritical()` and `ExitCritical()` as we used it does not allow nesting. Why?

Solution: The macros are using a global variable to store the status of the interrupts (CCR register). If nested, then that variable gets overwritten, causing wrong behavior of the application.

Question 126.....Points: [2]

Given the following source code:

```
#defineCALC1 (2+5)
#defineCALC2 (5*3)
```

Why is it important to use parenthesis?

Solution: The preprocessor does a textual replacement. Without the parenthesis it is possible that after the preprocessing the output is non-intentional. For example `10*CALC1` would result into `10*2+5`.

Question 127.....Points: [2]

List some advantages and disadvantages of using macros:

Solution: Pros: calculations at compilation time (faster, denser code), using it for configuration and compile time configuration. Cons: only textual replacement, overuse can make code harder to read, no type safety, debugging macros is a problem.

Question 128.....Points: [2]

You need to write a header file `LED.h`. Write it in such a way so *recursive* includes are not possible.

Solution:

```
#ifndef __LED_H_
#define __LED_H_
/* content of the header *.h file */
#endif
```

Question 129.....Points: [1]

You need to use a variable in multiple modules (.c files). Where do you put the *declaration* of it? Make an example.

Solution: The declaration needs to be in an interface header (*.h) file. E.g. `extern int MOT_currPWM;`. The name should be as such that there are no name collisions. Common practice is to use a prefix like `MOT_`.

Question 130.....Points: [1]

What is the fundamental difference between the two VCS: Git and SubVersion?

Solution: Git is a distributed VCS, while SubVersion is a centralized VCS.

Question 131.....Points: [1]

Two engineers are using Git as VCS. Now they change both the same line of code, and both perform a *commit*. What happens?

Solution: Nothing at the time of the commit. But at the time of *push*, the system will report a conflict which needs to be resolved manually.

Question 132.....Points: [3]

Below is the content of a .gitignore file. Explain what this file does with such a content:

```
*
!*.c
!*.h
!*.gitignore
!dev/
```

Solution: In the .gitignore the list of files and extensions tells Git which files to ignore. Git builds a list of files and goes from the beginning to the end of the file. The first line ignores all files. Then it removes (! operator) all *.c and *.h files, including the gitignore file. Then it removes all files in the dev folder. Using the ! operator gitignore explicitly specifies the files *not* to ignore.

Question 133.....Points: [3]

For multiple Eclipse projects, you want to use source files in a common folder (e.g. `c:/mySources/common`) which is *outside* of your Eclipse project. You consider to use 'Virtual Group', 'Linked Folder' or 'Linked File' methods to use the common file(s) in your project. Now you have added a new source file (e.g. `accelerometer.c`) to that common folder. For each of the three above methods, explain what you have to do for *each of your eclipse projects* which want to use the new `accelerometer.c` file.

Solution:

- Virtual Group: need to add another link in the group
- Linked Folder: nothing, it will be automatically added
- Linked Files: need to add another link

Question 134.....Points: [3]

Given a Quadrature encoder as in Figure 7, generating two signals C0 and C1. Both signals get sampled with a periodic timer interrupt. The disk is turning with a maximum frequency of 100 Hz. Determine the minimal sampling frequency in Hz to guarantee an error free sampling of the two signals C0 and C1.




Figure 7: Quad Disk

Solution: $(12 \cdot 4 \cdot 100 \text{ Hz} \cdot 2) = 9600 \text{ Hz}$.

1p: 12 sectors

1p: 4 signals each sector

1p: Shannon/Nyquist

Question 135.....Points: [3]

You are using a wireless transceiver as used in the lab to transmit sensor data values to another system. Describe strategies to reduce the energy consumption for the communication. Limit your answer to the *communication* and *transceiver* only.

Solution: Data packets as small as possible, sending as few as possible. Using sleep mode of the transceiver.

Question 136.....Points: [3]

Calculate for the 16bit **Binary Code 0x1722_b** the corresponding 16bit (binary reflected) Gray code_g:

Solution: Binary to Gray:

0x1722 EXOR (0x1722»1)

= 0x1722 EXOR 0x0B91 = **0x1CB3_b** (or **0b1'1100'1011'0011**)




Figure 8: 6bit Gray Code

Question 137.....Points: [3]

In Figure 8 are steps transforming a 6bit Gray Code into a 6bit Binary Code. Fill in the gaps:

Question 138.....Points: [8]

Below is the implementation of a task which reads and writes to an SCI (Serial Communication Interface). The SCI is using interrupts for RX and TX. As other tasks are using the `Read()` and `Write()` functions as well, it is necessary to use critical sections.

```
void Read(char *buf, size_t bufSize) {
    /* start critical section */
    SCI_Read(buf, bufSize);
    /* end critical section */
}

void Write(char *buf) {
    /* start critical section */
    SCI_Write(&buf[0]);
    /* end critical section */
}

static portTASK_FUNCTION(Task1, pvParameters) {
    unsigned char buf[16];
    (void)pvParameters; /* parameter not used */
    for(;;) {
        /* start critical section */
        Read(buf, sizeof(buf));
        Write(buf, sizeof(buf));
        /* end critical section */
    }
}
```

You identified 3 different ways for implementing the `/* start critical section`

`*/` and `/* end critical section */` in above source. Evaluate following statements for the consequences of each implementation for the above source:

- (a) With using `DisableAllInterrupts()` and `EnableAllInterrupts()` the following applies: [2]
- $\oplus \pm$ It increases the interrupt latency time.
 - $\oplus \pm$ Is the best solution with respect to RAM/ROM footprint.
 - $\oplus \pm$ Makes the usage of `SCI_Read()` and `SCI_Write()` reentrant.
 - $\ominus \pm$ Allows the scheduler to run inside every critical section.
- (b) With using FreeRTOS *binary semaphore/mutex* the following applies: [2]
- $\oplus \pm$ Calling the semaphore/mutex API might trigger a context switch.
 - $\ominus \pm$ No context switch will happen within the critical section.
 - $\oplus \pm$ Only one task at a time will be inside the critical section.
 - $\oplus \pm$ Interrupts will occur inside the critical section.
- (c) With using FreeRTOS *recursive semaphore/mutex* the following applies: [2]
- $\oplus \pm$ Calling the semaphore/mutex API might trigger a context switch.
 - $\ominus \pm$ Multiple tasks might be within the same critical section.
 - $\ominus \pm$ No context switch will happen within the critical section.
 - $\ominus \pm$ No interrupts will occur inside the critical section.
- (d) From the 3 solutions a), b) and c), select the one you consider as the best one for above source, and explain briefly why: [2]

Solution: recursive semaphore/mutex.

Question 139.....Points: [4]

A system using *Priority Ceiling* has four tasks J_1, J_2, J_3 and J_4 . The task index denotes the task priority with 4 as the highest priority. Additionally four semaphores S_1, S_2, S_3 and S_4 are used. Task are getting a lock of a semaphore x using LS_x , the lock of a semaphore is released again with US_x . The task and semaphores are used according to following schedule:

- $J_1 = \{LS_4, US_4\}$
- $J_2 = \{LS_2, LS_1, US_1, US_2, LS_3, US_3\}$
- $J_3 = \{LS_1, US_1, LS_3, US_3\}$
- $J_4 = \{LS_1, US_1\}$

Determine the *Priority Ceiling* for each semaphore:

Priority Ceiling for S_1 :

139. _____ 4 _____

Priority Ceiling for S_2 :

139. _____ 2 _____

Priority Ceiling for S_3 :

139. _____ **3** _____

Priority Ceiling for S_4 :

139. _____ **1** _____

Question 140.....Points: [3]

Below is a way to prevent spurious keyboard interrupts on a microcontroller:

```
PTADD = 0x00; /* set port as input */
PTAD = 0xFF; /* write 1's to data port to have defined input values */
PTAPE = 0xFF; /* enable pull ups */
EnableInterrupts;
```

Transform the above code into an implementation using *Gadfly* synchronization:

Solution:

```
PTADD = 0x00; /* set port as input */
PTAPE = 0xFF; /* enable pull ups */
while (PTAD!=0xFF); /* wait until port pull ups are working */
EnableInterrupts;
```

Question 141.....Points: [6]

A preemptive system has 3 tasks T_1 , T_2 and T_3 where the task index denotes the task priority, with 1 the lowest priority. The system is using *Priority Inheritance* for the semaphores. Create a timing diagram with the tasks in Figure 9 indicating the execution time for each task. The system has following timing:

- t_0 : T_1 and T_2 are ready, T_3 is running,
- t_1 : The running task gets suspended
- t_3 : The running task gets suspended
- t_4 : The running task requests the semaphore
- t_5 : T_2 and T_3 get ready
- t_7 : The running task requests the semaphore
- t_8 : The task having the semaphore releases the semaphore
- t_9 : The task having the semaphore releases the semaphore
- t_{10} : The running task terminates
- t_{11} : The running task terminates
- t_{12} : The running task terminates




Figure 9: Tasks

Question 142.....Points: [3]

Explain for what the *Anti-Windup* is used in a PID control system:

Solution: To limit the impact of the integral term in the PID.

Question 143.....Points: [3]

Identify the three serious problems in the following source:

```
#define QUEUE_LENGTH 5
#define QUEUE_ITEM_SIZE sizeof(char_t*)

void QUEUE_SendMessage(const char *msg) {
    char *ptr = FRTOSt_pvPortMalloc(UTIL1_strlen(msg));
    UTIL1_strcpy(ptr, msg);
    if (xQueueSendToBack(queueHandle, ptr, portMAX_DELAY) != pdPASS)
    {
        for (;;) {} /* ups? */
    }
}
```

Solution: It should be `UTIL1_strlen(msg)+1`, it does not check for `malloc()` failure, it does not pass the address of the pointer

Question 144 Points: [10]

FreeRTOS is used in priority based preemptive mode. Tick timer is set up for 5 ms, and the processor is running at maximum speed. Task T3 has priority 3 (highest priority), task T2 has priority 2 and task T1 has priority 1. The tasks have been created with `xTaskCreate()` before time $t=0$ ms, and the scheduler is started with `vTaskStartScheduler()` at the time $t=0$ ms. Draw in Figure 10 a timing diagram for the execution of the two tasks for the first 100 ms (after $t_0 = 0$ ms). Use a bar to indicate when a task is running. You can ignore the overhead in the task loop and the overhead in the RTOS/scheduler itself.

```
static portTASK_FUNCTION(T3, pvParameters) { /* priority 3 task */
    portTickType xLastTime = xTaskGetTickCount();
    for (;;) { /* task time is 7 ms including overhead */
        DoWorkFor7ms(); /* this needs 7 ms */
        vTaskDelayUntil(&xLastTime, 30/portTICK_RATE_MS);
    } /* loop forever */
}

static portTASK_FUNCTION(T2, pvParameters) { /* priority 2 task */
    for (;;) { /* task time is 7 ms including overhead */
        DoWorkFor7ms(); /* this takes 7 ms */
        vTaskDelay(30/portTICK_RATE_MS);
    } /* loop forever */
}

static portTASK_FUNCTION(T1, pvParameters) { /* priority 1 task */
    for (;;) { /* task time is 4 ms including overhead */
        DoWorkFor4ms(); /* this needs 4 ms */
        vTaskDelay(28/portTICK_RATE_MS);
    } /* loop forever */
}
```

}




Figure 10: Task Timing

Question 145.....Points: [3]

You are working with an Eclipse based project (you can consider the one we used for INTRO). Which kind of files are you going to put into a VCS? Which ones not? Explain briefly why.

Solution: Put all files into the repository which are needed to build the product/project (source files, project settings, configuration files), but do not put any derived or generated resources into it (generated documentation, object files, ELF file(s), generated make files, log files). The derived files can be re-generated, and putting them into VCS will just be duplication of content/consumes a lot of space.

Question 146 **Points (Bonus)**

Note: the points of the following section (max 34 Points) can be applied for the series of questions afterward (worth 30 points). The maximum number of points of each section counts.

- (a) Draw the path of the robot from the start vector to the start of the maze in Fig. 11. [5]
- (b) Draw the path of the robot in the maze using either the 'left-hand' or 'right-hand' rule in Fig. 11. [5]
- (c) Calculate the Challenge ranking point with the Swiss Rule, using the following formula: $26 - (Rank * 2)$. [24]

Solution:




Figure 11: Maze robot

Question 147 **Points: [1]**

The slide material for *this* semester week 2 are covering the following topics:

- ⊕ ± Systems and Realtime
- ⊖ ± GitHub and Version Control
- ⊕ ± Processor Expert and Drivers
- ⊕ ± Preprocessor and Macros

Question 148 **Points: [1]**

With $X == 0x00EFC000$ and $Y = ((X \gg 16) \& 0xFF) - 5$, what is Y in hexadecimal?

Solution: 0xEA

Question 149 **Points: [1]**

Systems can be classified as:

- ☐ converting, reacting, interacting
- ☐ measuring, controlling, calculating

☒ transforming, reactive, interactive

☐ interactive, responsive, collective

☐ soft realtime, hard realtime, true realtime

Question 150.....Points: [1]

What is the value of variable `a` after the following code:

```
int a = 0x1234;
a &= ~(1<<2);
```

Solution: 0x1230;

Question 151.....Points: [1]

Which function gets called in `foo()`?

```
#define CAT(a,b) a##b
void foo(void) {
    CAT(t,a)();
}
```

☐ `taab()`;

☐ `TaAb()`;

☒ `ta()`;

☐ `aTbA()`;

☐ `ab()`;

Question 152.....Points: [1]

What is the difference between `(size_t*)` and `(void*)`?:

Solution: The first is a pointer to the data type which is returned by the `sizeof()` operator. It denotes the needed type to hold a size type for the maximum size object in memory. For example if in a system it is only possible to have memory objects up to 64KByte, this would be of type `uint16_t`. The second type is a generic pointer type, and holds a pointer capable to point to any place in the memory. Both types are similar, but not the same.